

# Big Data Analytics Lab 2

*Omkar Bhutra (omkbh878) and Vinay Bengaluru (vinbe289)*

*17 May 2019*

## Execution of the code

```
[x_omkbh@heffa1 ~]$ ./runYarn.sh bda2alt.py
```

```
#Question 1
#year, station with the max, maxValue ORDER BY maxValue DESC
#year, station with the min, minValue ORDER BY minValue DESC

from pyspark import SparkContext
from pyspark.sql import SQLContext, Row
from pyspark.sql import functions as F

sc = SparkContext(appName = "Max_MinTemperatures_sparksql")
sqlContext = SQLContext(sc)
rdddata = sc.textFile('/user/x_omkbh/data/temperature-readings.csv') \
    .map(lambda line: line.split(";")) \
    .filter(lambda obs:
                (int(obs[1][:4]) >= 1950 and
                 int(obs[1][:4]) <= 2014)) \
    .map(lambda obs: \
        Row(station = obs[0], date = obs[1], \
            year = obs[1].split("-")[0], time = obs[2], \
            value = float(obs[3]), quality = obs[4]))
schemaTempReadings = sqlContext.createDataFrame(rdddata)
schemaTempReadings.registerTempTable("schemaTempReadings")

#Q1.1 year, station with the max, maxValue ORDER BY maxValue DESC

maxTemp = sqlContext.sql"""
    SELECT DISTINCT(table1.year) AS year,
           FIRST(table1.station) AS station,
           FIRST(value) AS value
    FROM schemaTempReadings AS table1
    INNER JOIN
    (
    SELECT year, MAX(value) AS max_value
    FROM schemaTempReadings
    GROUP BY year
    ) AS table2
    ON table1.year = table2.year
    WHERE table1.value = table2.max_value
    GROUP BY table1.year
    ORDER BY value DESC
    """
maxTemp.rdd.repartition(1) \
    .sortBy(ascending=False, keyfunc=lambda (year,station, value): value)
```

```

maxTemp.take(10)
maxTemp.write.save('/user/x_omkbh/bda2.11')

#Q1.2 year, station with the min, minValue ORDER BY minValue DESC

minTemp = sqlContext.sql("""
    SELECT DISTINCT(table1.year) AS year,
           FIRST(table1.station) AS station,
           FIRST(value) AS value
    FROM schemaTempReadings AS table1
    INNER JOIN
    (
        SELECT year, MIN(value) AS min_value
        FROM schemaTempReadings
        GROUP BY year
    ) AS table2
    ON table1.year = table2.year
    WHERE table1.value = table2.min_value
    GROUP BY table1.year
    ORDER BY value DESC
    """)
minTemp.rdd.repartition(1) \
    .sortBy(ascending=False, keyfunc=lambda (year,station, value): value)
minTemp.take(10)
minTemp.write.save('/user/x_omkbh/bda2.12')

```

**Output: Max Temp**

```

maxTemp.take(10)
[Row(year=u'1975', station=u'86200', value=36.1),
 Row(year=u'1992', station=u'63600', value=35.4),
 Row(year=u'1994', station=u'117160', value=34.7),
 Row(year=u'2014', station=u'96560', value=34.4),
 Row(year=u'2010', station=u'75250', value=34.4),
 Row(year=u'1989', station=u'63050', value=33.9),
 Row(year=u'1982', station=u'94050', value=33.8),
 Row(year=u'1968', station=u'137100', value=33.7),
 Row(year=u'1966', station=u'151640', value=33.5),
 Row(year=u'2002', station=u'78290', value=33.3)]

```

**Output: Min Temp**

```

minTemp.take(10)
[Row(year=u'1990', station=u'147270', value=-35.0),
 Row(year=u'1952', station=u'192830', value=-35.5),
 Row(year=u'1974', station=u'166870', value=-35.6),
 Row(year=u'1954', station=u'113410', value=-36.0),
 Row(year=u'1992', station=u'179960', value=-36.1),

```

```

Row(year=u'1975', station=u'157860', value=-37.0),
Row(year=u'1972', station=u'167860', value=-37.5),
Row(year=u'1995', station=u'182910', value=-37.6),
Row(year=u'2000', station=u'169860', value=-37.6),
Row(year=u'1957', station=u'159970', value=-37.8)]

```

*#Question 2:*

*#year, month, value ORDER BY value DESC*

*#year, month, value ORDER BY value DESC*

```

rdddata2 = sc.textFile('/user/x_omkbh/data/temperature-readings.csv') \
    .map(lambda line: line.split(";")) \
    .map(lambda obs: \
        Row(station = obs[0], date = obs[1], \
            year = obs[1].split("-")[0], month = obs[1].split("-")[1], \
            yymm = obs[1][:7], \
            time = obs[2], value = float(obs[3]), quality = obs[4]))
schemaTempReadings2 = sqlContext.createDataFrame(rdddata2)
schemaTempReadings2.registerTempTable("schemaTempReadings2")

```

*#Q2.1 Temperatures readings higher than 10 degrees*

```

overTenTemp = sqlContext.sql(" \
    SELECT FIRST(year), FIRST(month), COUNT(value) AS counts\
    FROM schemaTempReadings2 \
    WHERE value >= 10 AND year >= 1950 AND year <= 2014\
    GROUP BY year, month \
    ORDER BY counts DESC")

```

*#Q2.2 Distinct Temperatures readings higher than 10 degrees*

```

overTenTempDistinct = schemaTempReadings2.filter(schemaTempReadings2["value"] > 10) \
    .groupBy("yymm") \
    .agg(F.countDistinct("station").alias("count"))
overTenTempDistinct = overTenTempDistinct.rdd.repartition(1) \
    .sortBy(ascending = False, keyfunc = lambda \
        (yymm, counts): counts)
overTenTempDistinct.saveAsTextFile('/user/x_omkbh/bda2.2')

```

Output of Distinct Temperatures readings counts:

```

print overTenTempDistinct.take(10)
[Row(yymm=u'1972-10', count=378),
 Row(yymm=u'1973-05', count=377),
 Row(yymm=u'1973-06', count=377),
 Row(yymm=u'1973-09', count=376),
 Row(yymm=u'1972-08', count=376),
 Row(yymm=u'1972-05', count=375),
 Row(yymm=u'1972-06', count=375),
 Row(yymm=u'1972-09', count=375),

```

```
Row(yymm=u'1971-08', count=375),
Row(yymm=u'1972-07', count=374)]
```

*#Question 3.*

*#year, month, station, avgMonthlyTemperature ORDER BY avgMonthlyTemperature DESC*

```
rdddata3 = sc.textFile('/user/x_omkbh/data/temperature-readings.csv') \
    .map(lambda line: line.split(";")) \
    .filter(lambda p:
        (int(p[1][:4]) >= 1950 and
         int(p[1][:4]) <= 2014)) \
    .map(lambda p: Row(station=int(p[0]),
        day=p[1],
        month=p[1][:7],
        value=float(p[3])))
schemaTempReadings3 = sqlContext.createDataFrame(rdddata3)
schemaTempReadings3.registerTempTable("schemaTempReadings3")
avgMonthTemp = sqlContext.sql(
    """
    SELECT mytbl.month, mytbl.station, AVG(mytbl.max_value + mytbl.min_value) / 2 AS avg_value
    FROM
    (
    SELECT month, station, MIN(value) AS min_value, MAX(value) AS max_value
    FROM schemaTempReadings3
    GROUP BY day, month, station
    ) AS mytbl
    GROUP BY mytbl.month, mytbl.station
    ORDER BY AVG(mytbl.max_value + mytbl.min_value) / 2 DESC
    """
)
avgMonthTemp.rdd.repartition(1).sortBy(ascending=False,
    keyfunc=lambda (month, station, value): value)
avgMonthTemp.write.save('/user/x_omkbh/bda2.3')
```

Output Average monthly temperatures:

```
print avgMonthTemp.take(10)
[Row(month=u'2014-07', station=96000, avg_value=26.3),
Row(month=u'1994-07', station=96550, avg_value=23.07105263157895),
Row(month=u'1983-08', station=54550, avg_value=23.0),
Row(month=u'1994-07', station=78140, avg_value=22.970967741935485),
Row(month=u'1994-07', station=85280, avg_value=22.872580645161293),
Row(month=u'1994-07', station=75120, avg_value=22.858064516129033),
Row(month=u'1994-07', station=65450, avg_value=22.856451612903232),
Row(month=u'1994-07', station=96000, avg_value=22.808064516129033),
Row(month=u'1994-07', station=95160, avg_value=22.764516129032256),
Row(month=u'1994-07', station=86200, avg_value=22.711290322580645)]
```

*#Question 4.*

*#station, maxTemp, maxDailyPrecipitation ORDER BY station DESC*

*#Note: The correct result for this question should be empty.*

```

# Temperatures
temperature_data = sc.textFile('/user/x_omkbh/data/temperature-readings.csv')
temperature_obs = temperature_data.map(lambda line: line.split(";")) \
    .map(lambda obs: Row(station=int(obs[0]),
                        temp=float(obs[3])))

schema_temp_readings = sqlContext.createDataFrame(temperature_obs)
schema_temp_readings.registerTempTable("temp_readings")

# precipitation
precipitation_data = sc.textFile('/user/x_omkbh/data/precipitation-readings.csv')
precipitation_obs = precipitation_data.map(lambda line: line.split(";")) \
    .map(lambda obs: Row(station=int(obs[0]),
                        day=obs[1],
                        precip=float(obs[3])))

schema_precip_readings = sqlContext.createDataFrame(precipitation_obs)
schema_precip_readings.registerTempTable("precip_readings")

combined = sqlContext.sql(
    """
    SELECT tr.station, MAX(temp) AS max_temp, MAX(precip) AS max_precip
    FROM schema_temp_readings AS tr
    INNER JOIN
    (
    SELECT station, SUM(precip) AS precip
    FROM schema_precip_readings
    GROUP BY day, station
    ) AS pr
    ON tr.station = pr.station
    WHERE (temp >= 25 AND temp <= 30)
    AND (precip >= 100 AND precip <= 200)
    GROUP BY tr.station
    ORDER BY tr.station DESC
    """
)

tempPrec = combined.rdd.repartition(1) \
    .sortBy(ascending=False, keyfunc=lambda (station, temp, precip): station)

tempPrec.take(10)
tempPrec.saveAsTextFile('/user/x_omkbh/bda2.4')

```

Output Max daily temperatures/precipitation:

```

#Station number, maximum measured temperature, maximum daily precipitation
(u'128510', (29.5, None))
(u'192830', (29.5, None))
(u'84660', (27.6, None))
(u'139110', (29.0, None))
(u'161670', (25.7, None))
(u'166940', (27.9, None))
(u'77180', (29.3, None))
(u'180740', (29.0, None))
(u'72340', (29.8, None))

```

```
(u'147560', (29.9, None))
(u'180750', (29.3, None))
(u'83460', (28.0, None))
(u'83620', (29.4, None))
(u'159680', (26.2, None))
(u'139340', (28.9, None))
```

*#Question 5.*

*#year, month, avgMonthlyPrecipitation ORDER BY year DESC, month DESC*

```
ostergotlandStations = sc.textFile('/user/x_omkbh/data/stations-Ostergotland.csv') \
    .map(lambda line: line.split(";")) \
    .map(lambda obs: int(obs[0])) \
    .distinct().collect()
ostergotlandStations = sc.broadcast(ostergotlandStations)
ostergotlandStations = {station: True for station in ostergotlandStations.value}
precipitations = sc.textFile('/user/x_omkbh/data/precipitation-readings.csv') \
    .map(lambda line: line.split(";")) \
    .filter(lambda obs:
        ostergotlandStations.get(int(obs[0]), False)) \
    .map(lambda obs: Row(day=obs[1],
        month=obs[1][:7],
        station=int(obs[0]),
        precip=float(obs[3])))

precSchema = sqlContext.createDataFrame(precipitations)
precSchema.registerTempTable("PrecSchema")
avgMthPrec = sqlContext.sql(
    """
    SELECT mytbl2.month, AVG(mytbl2.precip) AS avg_precip
    FROM
    (
    SELECT mytbl1.month, mytbl1.station, SUM(mytbl1.precip) AS precip
    FROM
    (
    SELECT month, station, SUM(precip) AS precip
    FROM PrecSchema
    GROUP BY day, month, station
    ) AS mytbl1
    GROUP BY mytbl1.month, mytbl1.station
    ) AS mytbl2
    GROUP BY mytbl2.month
    ORDER BY mytbl2.month DESC
    """
)
avgMthPrec.rdd.repartition(1).sortBy(ascending=False, keyfunc=lambda (month, precip): month)
avgMthPrec.saveAsTextFile('/user/x_omkbh/bda2.5')
```

Output Ostergotland average monthly precipitation:

```
print avgMthPrec.take(10)
[Row(month=u'2016-07', avg_precip=0.0),
 Row(month=u'2016-06', avg_precip=47.6625),
```

```

Row(month=u'2016-05', avg_precip=29.250000000000004),
Row(month=u'2016-04', avg_precip=26.9),
Row(month=u'2016-03', avg_precip=19.9625),
Row(month=u'2016-02', avg_precip=21.5625),
Row(month=u'2016-01', avg_precip=22.325),
Row(month=u'2015-12', avg_precip=28.925),
Row(month=u'2015-11', avg_precip=63.887499999999996),
Row(month=u'2015-10', avg_precip=2.2625)]

```

*#Question 6.*

*#year, month, difference ORDER BY year DESC, month DESC*

*# Ostergotland Stations*

```

ostergotlandStations = sc.textFile('/user/x_omkbh/data/stations-Ostergotland.csv') \
    .map(lambda line: line.split(";")) \
    .map(lambda obs: int(obs[0])) \
    .distinct().collect()
ostergotlandStations = sc.broadcast(ostergotlandStations)
ostergotlandStations = {station: True for station in ostergotlandStations.value}
temperatures = sc.textFile('/user/x_omkbh/data/temperature-readings.csv') \
    .map(lambda line: line.split(";")) \
    .filter(lambda obs: ostergotlandStations.get(int(obs[0]), False)) \
    .map(lambda obs: \
        Row(station = obs[0], \
            date = obs[1], \
            year = obs[1].split("-")[0], \
            month = obs[1].split("-")[1], \
            day = obs[1].split("-")[2], \
            yymm = obs[1][:7], \
            yymmdd = obs[1], \
            time = obs[2], \
            temp = float(obs[3]), \
            quality = obs[4]))
tempSchema = sqlContext.createDataFrame(temperatures)
tempSchema.registerTempTable("TempSchema")
avgMthTemp = sqlContext.sql("""
    SELECT one.yymm,
           AVG(one.minTemp + one.maxTemp) / 2 AS avgTemp
    FROM
    (
    SELECT yymm,
           year,
           yymmdd,
           MIN(temp) AS minTemp,
           MAX(temp) AS maxTemp
    FROM TempSchema
    GROUP BY yymmdd,
             yymm,
             year,
             station
    ) AS one
    WHERE one.year >= 1950 AND one.year <= 2014
    GROUP BY one.yymm
    """)

```

```

longTermAvgTemp = avgMthTemp.filter(avgMthTemp.substring(avgMthTemp["yymm"], 1, 4) <= 1980) \
    .groupBy(avgMthTemp.substring(avgMthTemp["yymm"], 6, 7).alias("month")) \
    .agg(avgMthTemp.avg(avgMthTemp["avgTemp"]).alias("longTermAvgTemp"))
diffTemp = avgMthTemp.join(longTermAvgTemp,
    (avgMthTemp.substring(avgMthTemp["yymm"], 6, 7) ==
    longTermAvgTemp["month"]), "inner")
diffTemp = diffTemp.select(diffTemp["yymm"],
    (diffTemp.abs(diffTemp["avgTemp"]) -
    diffTemp.abs(diffTemp["longTermAvgTemp"])).alias("diffTemp"))
diffTemp = diffTemp.rdd.repartition(1).sortBy(ascending = False,
    keyfunc = lambda (yymm, diff): yymm)
diffTemp.write.save('/user/x_omkbh/bda2.6')

```

Output Ostergotland average monthly precipitation temperature difference:

```

#take(13)
1950,01,2.00483133412
1950,02,-2.34798988599
1950,03,1.1819828212
1950,04,1.60069315899
1950,05,0.982351940463
1950,06,-0.216232256095
1950,07,-1.47714267742
1950,08,0.241517150903
1950,09,0.343179398558
1950,10,-0.460520515247
1950,11,-0.47779366064
1950,12,1.07259158462
1951,01,-0.19629769814
1951,02,-2.60656131457
1951,03,3.08359572443
1951,04,-0.0381401743418
1951,05,-1.93038999502

```

```

part.000006 <- read.csv("C:/Users/Omkar/Downloads/Big Data Analytics/lab2/part-000006", header=FALSE)
part.000006 %>% ggplot(aes(x=V1, y=V3))+geom_point()+
  ggtitle("Average temperature difference")+
  labs(x="Year",y="average temperature difference")

```



